

STAT 527: Assignment #2

Philip Pham

May 1, 2020

Problem 1

(a) *Proof.* The likelihood function is the probability of observing the data.

$$\begin{aligned} L(\{y_{i,n}\}; f, \sigma) &= \prod_{i=1}^n P(y_{i,n}; f, \sigma) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2} \left(\frac{y_{i,n} - f(i/n)}{\sigma}\right)^2\right) \end{aligned}$$

□

(b) The approach is reasonable especially if we choose f such that the likelihood is a differentiable function of the parameters. Then, it is easy to choose the optimal \hat{f} .

Such an approach may be undesirable if our family of functions is too simple (high bias) or is not expressive enough (high variance). There may be other issues like the number of parameters that can lead to overfitting.

(c) We can maximize the log-likelihood:

$$l(\{y_{i,n}\}; \mu, \sigma) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_{i,n} - \mu)^2.$$

Taking the derivative with respect to μ and solving, we get $\hat{\mu} = n^{-1} \sum_{i=1}^n y_{i,n}$, that is, the mean of observations.

For the localization, we take a local mean:

$$\begin{aligned} f(x_0) &= \frac{\sum_{i=1}^n K_h(1/n - x_0) y_{i,n}}{\sum_{i=1}^n K_h(1/n - x_0)}, \\ K_h(x) &= \begin{cases} 1, & |x| \leq h; \\ 0, & \text{otherwise.} \end{cases} \end{aligned}$$

This is just the Nadaraya-Watson estimator with a box kernel with bandwidth h .

Problem 2

Proof. Note that

$$\hat{\mu}_n = \frac{1}{n} \sum_{i \leq n} x_i \sim \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right)$$

$$n \frac{\hat{\sigma}_n^2}{\sigma^2} \sim \chi_{n-1}^2 \Rightarrow \hat{\sigma}_n^2 = \frac{1}{n} \sum_{i \leq n} (x_i - \hat{\mu})^2 \sim \frac{\sigma^2}{n} \chi_{n-1}^2.$$

Thus, we have $E[\hat{\sigma}_n^2] = \frac{n-1}{n} \sigma^2$ and $\text{var}(\hat{\sigma}_n^2) = 2 \frac{n-1}{n^2} = 2/n - 2/n^2$. Thus, the estimators are consistent. Using independence, Slutsky's theorem, and central limit theorem, we'll have that

$$\sqrt{n} \left(\begin{pmatrix} \hat{\mu}_n \\ \hat{\sigma}_n \end{pmatrix} - \begin{pmatrix} \mu \\ \sigma \end{pmatrix} \right) \xrightarrow{D} \mathcal{N}\left(\mathbf{0}, \begin{pmatrix} \sigma^2 & 0 \\ 0 & 2 \end{pmatrix}\right).$$

The delta method tells us that

$$\sqrt{n} (\phi_{\hat{\mu}, \hat{\sigma}}(x_0) - \phi_{\mu, \sigma}(x_0)) \xrightarrow{D} \mathcal{N}\left(\mathbf{0}, \nabla \phi_{\mu, \sigma}(x_0)^\top \begin{pmatrix} \sigma^2 & 0 \\ 0 & 2 \end{pmatrix} \nabla \phi_{\mu, \sigma}(x_0)\right),$$

where the gradient is evaluated with respect to μ and σ .

This just multiplies the diagonals by some constants and doesn't change the rate of convergence. Convergence in distribution implies stochastic boundness, so squaring both sides we have that

$$(\phi_{\hat{\mu}, \hat{\sigma}}(x_0) - \phi_{\mu, \sigma}(x_0))^2 = O_p(1/n).$$

□

Problem 3

(a) See Figures 1, 2, and 3

While the trend isn't perfect, in general the larger k tend to choose smaller bandwidths and more closely fit the data or sometimes noise.

(b) After running 100 random folds, the following bandwidths were found.

```
[Summary(f='sin', kernel='box', oracle=0.25, k=2, mean=0.42499999999999993, var=0.025345),
Summary(f='sin', kernel='box', oracle=0.25, k=5, mean=0.41139999999999993, var=0.03555404),
Summary(f='sin', kernel='box', oracle=0.25, k=10, mean=0.39749999999999985, var=0.040248750000000014),
Summary(f='sin', kernel='gaussian', oracle=0.13, k=2, mean=0.20759999999999998, var=0.01192024),
Summary(f='sin', kernel='gaussian', oracle=0.13, k=5, mean=0.14210000000000003, var=0.0072085900000000026),
Summary(f='sin', kernel='gaussian', oracle=0.13, k=10, mean=0.12510000000000002, var=0.0066629899999999999),
Summary(f='poly', kernel='box', oracle=0.11, k=2, mean=0.11530000000000001, var=0.0013089100000000008),
Summary(f='poly', kernel='box', oracle=0.11, k=5, mean=0.08289999999999997, var=0.00011058999999999996),
Summary(f='poly', kernel='box', oracle=0.11, k=10, mean=0.08079999999999998, var=1.5359999999999999e-05),
Summary(f='poly', kernel='gaussian', oracle=0.08, k=2, mean=0.061200000000000004, var=0.00032255999999999984),
Summary(f='poly', kernel='gaussian', oracle=0.08, k=5, mean=0.049699999999999994, var=6.491000000000001e-05),
Summary(f='poly', kernel='gaussian', oracle=0.08, k=10, mean=0.04869999999999999, var=1.731000000000001e-05),
Summary(f='inv', kernel='box', oracle=0.33, k=2, mean=0.4341, var=0.03649419),
Summary(f='inv', kernel='box', oracle=0.33, k=5, mean=0.36940000000000006, var=0.028043640000000015),
Summary(f='inv', kernel='box', oracle=0.33, k=10, mean=0.42399999999999993, var=0.024012000000000002),
Summary(f='inv', kernel='gaussian', oracle=0.2, k=2, mean=0.3048, var=0.06769296),
Summary(f='inv', kernel='gaussian', oracle=0.2, k=5, mean=0.2088, var=0.014752560000000001),
Summary(f='inv', kernel='gaussian', oracle=0.2, k=10, mean=0.1886, var=0.008036040000000001)]
```

In general, bigger k chose smaller bandwidths with less variance. There are some exceptions like with $k = 10$ and $f(x) = (1 + 25x^2)^{-1}$. This tendency to choose a smaller bandwidth can sometimes lead to overfitting as it chooses a bandwidth smaller than the the oracle.

Figure 1: 2-fold cross validation

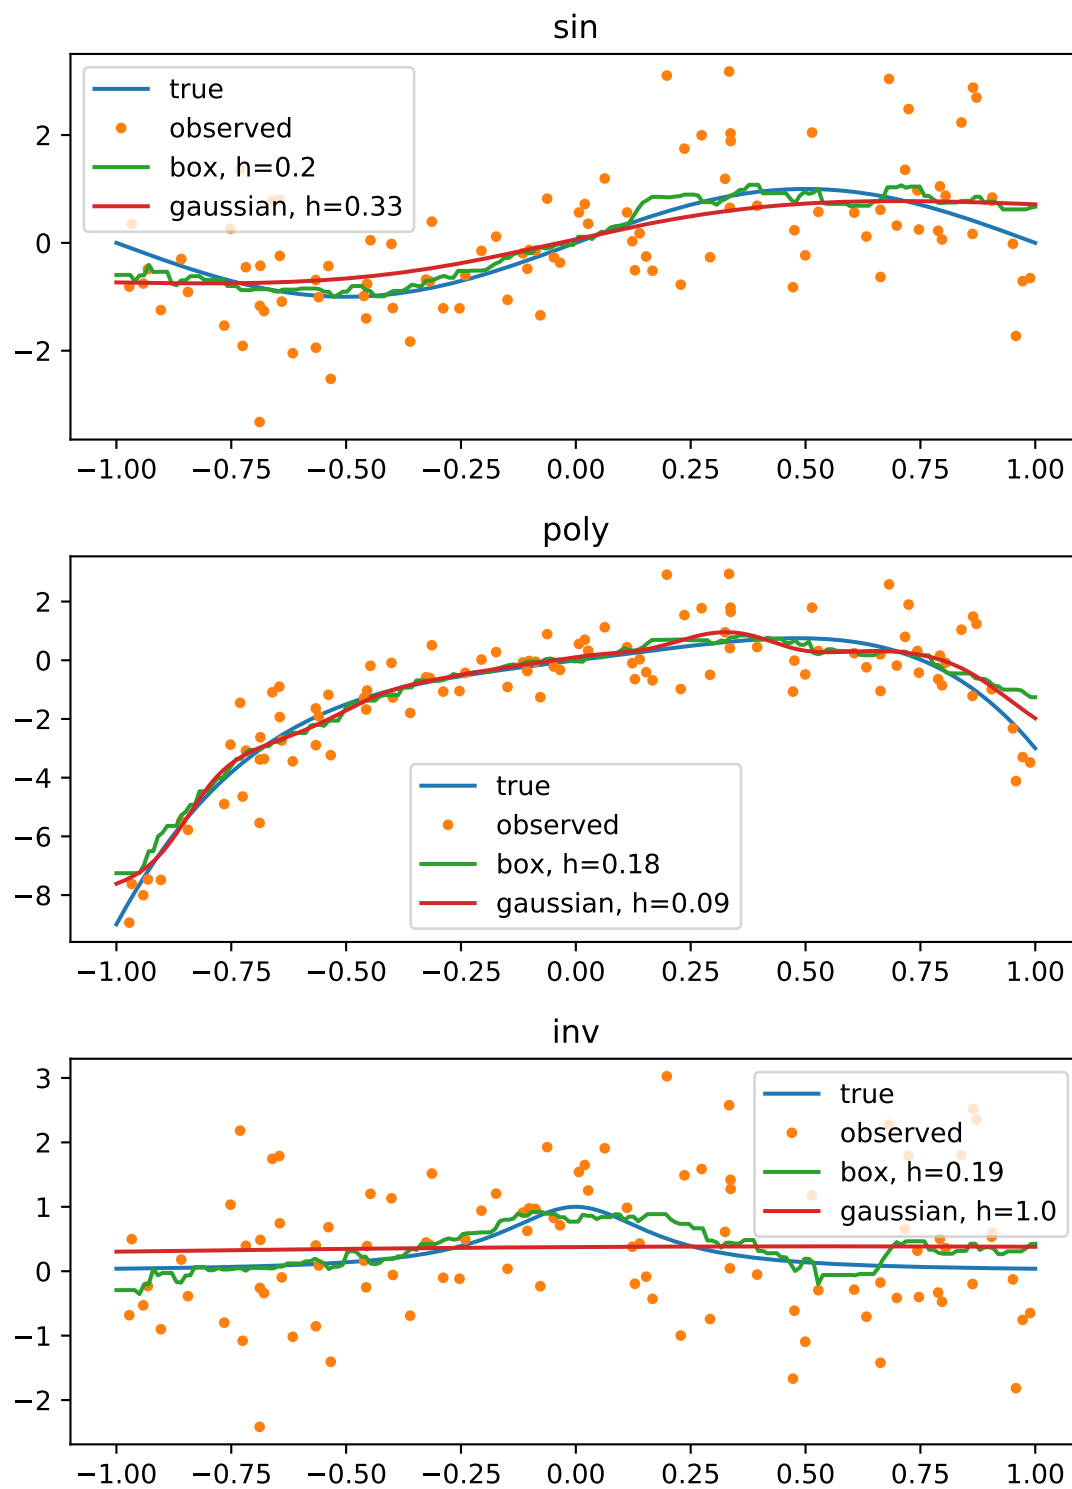


Figure 2: 5-fold cross validation

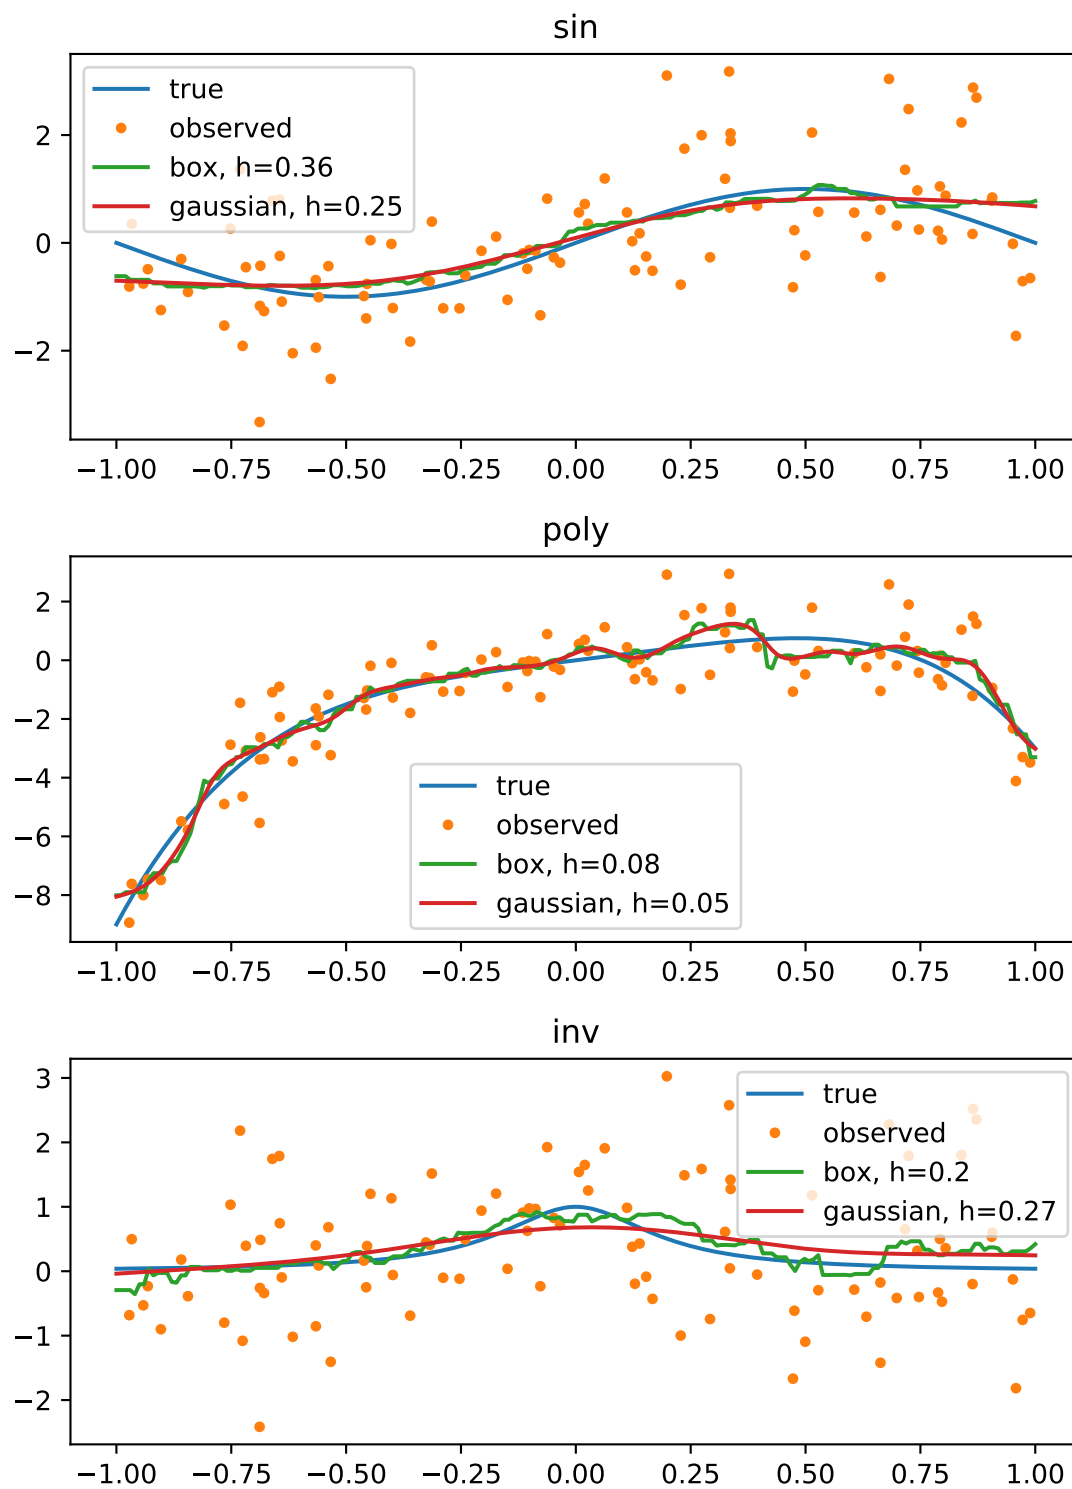


Figure 3: 10-fold cross validation

